

COEN 145 : Parallel Computing

Lab 3 : OPENMP

Part 0 : Request HPC resources

```
#!/bin/bash
```

```
# Syntax : #SBATCH --parameter=value
```

```
#SBATCH --job-name=pl-lab3
```

```
#SBATCH --partition=cpu
```

```
# CPU Cores = 2*omp_get_max_threads()
```

```
#SBATCH --cpus-per-task=X
```

```
#SBATCH --mem-per-cpu=8G
```

```
#SBATCH --nodes=1
```

```
#SBATCH --output=pl-lab3-%j.out
```

```
#SBATCH --time=10:00
```

```
#SBATCH --mail-type=ALL
```

```
#SBATCH --mail-user=USERNAME@scu.edu
```

```
# Environment Variables
```

```
export OMP_NUM_THREADS=X
```

```
export OMP_PLACES=cores
```

```
export OMP_PROC_BIND=true
```

```
# Build and Run
```

```
make clean && make
```

```
./mat_vec_mult.out > out.txt
```

```
./mat_vec_mult.out in_mat.txt in_vec.txt 1000 5000 5000 > out.txt
```

- Use the provided **run.sh** shell script to run **sum.out** from last week. Also do it in an **interactive** session

```
$ sbatch run.sh
```

```
$ srun --nodes=1 --ntasks=1 --cpus-per-task=2 --mem=8G --pty /bin/bash &&  
./sum.out 0 1 2 3 4 5 6 7 8 9
```

```
> sum = 45
```

```
> time-serial = X microseconds
```

Part 1 : Thread (Affinity/Pinning/Binding)

```
# specify number of OMP threads
```

```
$ export OMP_NUM_THREADS=2
```

```
# enable thread binding and print out info on thread affinity
```

```
$ export OMP_PROC_BIND=true
```

```
# bind each thread to a core
```

```
$ export OMP_PLACES=cores
```

Part 2 : Linking library inside Makefile



- Update our Makefile to include OPENMP using the compiler flag **-fopenmp**
- Print the **num_threads** available on the working machine using the functions `omp_get_num_threads()` and `omp_get_max_threads()`

```
> num_threads = X
```

```
> max_threads = Y
```

Part 3 : Measuring Performance

$$speedup = \frac{execution-time_{serial}}{execution-time_{parallel}}$$



Part 4 : Parallelize *sum.c*

- Using the decorators discussed in class modify your code to perform the summation reduction in parallel and report the output in a log file.

```
> sum = 45
```

```
> time-serial = X microseconds
```

```
> time-parallel = Y microseconds
```

```
> speedup = Z.xx (up to 2 decimal places)
```

Part 5 : Parallelize *mat_vec_mult.c*

- Write a program to perform a matrix vector multiplication in C (not C++) from randomly initialized values as in lab1. Also it should accept optional command-line arguments and read data from a text file as in lab1 with each element on a newline
- Implement using naive method with 2D array
- Compare performance against a single 1D buffer
- Parallelize your code using the directives covered in class and redirect (>) the output of the calculation to **out.txt** and create **figures** for speedup

```
$ ./mat_vec_mult.out {in_mat.txt | in_vec.txt | nrow | ncol1 | ncol2} > out.txt
```

```
> time-serial (2D) = X microseconds
```

```
> time-parallel (2D) = XZ microseconds
```

```
> time-serial (1D) = Y microseconds
```

```
> time-parallel (1D) = YZ microseconds
```

```
> speedup (2D) = A.xx (up to 2 decimal places)
```

```
> speedup (1D) = B.xx (up to 2 decimal places)
```