

Sparse matrix multiplication (OpenMP)

Introduction

The purpose of this assignment is to become familiar with OpenMP by implementing sparse matrix-matrix multiplication.

Sparse matrix multiplication ($C = AB$)

You need to write a program, *sparsematmult*, that will take as input two matrices A and B , and will output their product. A stub for your program is provided to you as C++ code. It is just an example. You can write your OpenMP code in C or C++. Complete the program by writing the *matmult* function. Your program should be able to be executed as:

```
sparsematmult 4 5 3 0.1 [optional parameters]
```

which would create two sparse matrices with random values, A of size 4×5 and B of size 5×3 , each with a fill factor of 0.1, and would multiply them, storing the result into a sparse matrix C of size 4×3 . Please note that the CSR of matrices A and B are given. Depending on the algorithm used for multiplication, you may need to transpose one of the matrices before multiplying them. The program should print out execution time in microseconds with up to 4 decimal points. You may optionally add a 4th parameter for *nthreads* (number of threads the program should execute the multiplication with).

General program

Design your programs so that they follow the following steps:

1. A single thread creates the matrices.
2. Multiple threads work together to do the matrix multiplication.
3. At the end, one thread reports the execution time.

These are just guidelines, and you should think carefully about step 2. Also, make sure the matrix multiplication computation is correct.

Parallelization strategy

Think about the amount of work that each thread does, given the sparsity patterns of the matrices.

Testing

Test your code on the HPC compute nodes (the *cpu* or *mem* queues) with the following parameters. For each execution, store results in log file which you can later use to create tables or figures for your report.

Fill factor: 0.05, 0.10, 0.15, 0.20

Number of threads: 1, 2, 4, 8, 12, 14, 16, 20, 24, 28.

Matrices:

A n rows	A n cols	B n cols
10000	10000	10000
20000	5300	50000
9000	35000	5750

Note that the evaluation of your code may be done using a different set of matrices than the ones above. As such, do not over-specify your code to work perfectly on those matrices and less so on others not in that set.

What you need to turn in

1. The source code of your program.
2. A short report that includes the following:
 - a. A short description of how you went about parallelizing sparse matrix multiplication. One thing that you should be sure to include is how you decided what work each thread would be responsible for doing.
 - b. Timing results for your parallel executions. These results should be reported in a strong scaling chart reporting the speedup at different numbers of threads vs the serial execution (1 thread).
 - c. A brief analysis of your results. Some things to consider might be:
 - i. How does the number of threads affect the runtime? Why do you think that is?
 - ii. How does the size of the problem affect the runtime? Why do you think that is?
 - iii. How well does your program scale, i.e., if you keep increasing the number of threads, do you think the performance will keep increasing at the same rate?
 - iv. Are performance characteristics consistent for the different shapes of the matrices? Why do you think that is?

Submission specifications

- A makefile must be provided to compile and generate the executable file.
- The executable file should be named 'sparsematmult'.
- All files (code + report) MUST be in a single directory and the directory's name MUST be your university student ID. Your submission directory MUST include at least the following files (other auxiliary files may also be included):

```
<Student ID>/sparsematmult.cpp  
<Student ID>/Makefile  
<Student ID>/report.pdf
```

- Submission MUST be in .tar.gz
- The following sequence of commands should work on your submission file:

```
tar xzvf <Student ID>.tar.gz  
cd <Student ID>  
make
```

```
ls -ld sparsematmult
```

This ensures that your submission is packaged correctly, your directory is named correctly, your makefile works correctly, and your output executable files are named correctly. If any of these does not work, modify it so that you do not lose points. I can answer questions about correctly formatting your submission BEFORE the assignment is due. Do not expect questions to be answered the night it is due.

Evaluation criteria

The goal for this assignment is for you to become familiar with the OpenMP API and develop an efficient parallel program. As such, the following things will be evaluated:

1. follows the assignment directions,
2. solve the problem correctly,
3. do so in parallel,
4. achieve speedup:
 - 5 / 10 points will be reserved for this criterion.